Timing Plan + Research Agent

Total time: 2 minutes 30 seconds -- one speaker, live demo

Segment	Clock	Time	Focus
1. Opening	0:00 - 0:10	10 sec	Set up the two agents
2. Research Agent	**0:10 - 1:15**	**65 sec**	**Cited answers**
3. Benchmark Agent	**1:15 - 2:20**	**65 sec**	**Safe comparison**
4. Wrap-up	2:20 - 2:30	10 sec	Closing line

0:00-0:10

1. OPENING

SHOW: Have the app open, logged in, with at least one uploaded filing (e.g. Apple FY2024 demo PDF).

SAY:

Now I will show two on-demand agents: the Research Agent, which answers questions with citations you can verify, and the Benchmark Agent, which compares companies without ever inventing a number.

0:10-1:15

2. RESEARCH AGENT (65 sec)

SHOW: Open the Research / chat page. Type the question: 'What was the total revenue in FY2024 and what are the key risks mentioned in the filing?' Then send. When the answer appears, point at the [S1] [S2] tags. Open one citation to show the exact source passage (document name, section, excerpt).

SAY:

The Research Agent is NOT allowed to answer from memory. First it retrieves real passages from the uploaded filings -- by meaning when embeddings are on, by keywords otherwise. Then the AI may answer ONLY from those passages, and every claim carries a tag like [S1] that points to the exact source. Watch -- if I ask something the documents do not contain, it says 'no source passages passed the relevance filter' instead of inventing an answer. This is retrieval-grounded research: every sentence is auditable.

Benchmark Agent + Wrap-up + Backup

1:15-2:20

3. BENCHMARK AGENT (65 sec)

SHOW: Open the Benchmark page. Select Apple vs Microsoft for the same fiscal year (e.g. FY2024). Generate the comparison. Point at: (a) the side-by-side metrics table, (b) the period labels under each company, (c) the verdict/insight text. Then mention the guard: pick different fiscal years and show it warns instead of comparing blindly.

SAY:

The Benchmark Agent compares companies safely -- three hard rules built into the code. Rule one: periods must be compatible -- it refuses to silently compare FY2024 with FY2023 and shows a clear warning instead. Rule two: the exact extracted metric values are

the

alyst

2:20-2:30

4. WRAP-UP (10 sec)

SHOW: Stay on the benchmark result screen.

SAY:

Cited research and safe comparisons -- two agents, zero invented numbers, and every figure traceable back to the filing. Thank you.

IF THE EXAMINER ASKS ABOUT THE CODE -- quick anchors

`src/lib/agents/researchAgent.ts -> answerResearchQuestion()`

Decomposes the question into retrieval steps, builds citations S1, S2... (document, section, chunk, excerpt), and answers only from them.

`src/lib/agents/benchmarkAgent.ts -> generateBenchmarkInsights()`

Period-compatibility guard, exact metric values packed by code, verdict compiled only from verified fields.

`src/lib/agents/documentAgent.ts -> searchDocumentChunks() + embeddingAgent.ts (optional vectors)`

Q&A BACKUP (one-line answers)

Q: What if the filing has no relevant passage?
A: It answers 'no source passages passed the relevance filter' -- never invents.

Q: What if embeddings are not configured?
A: Keyword relevance search takes over automatically -- citations work the same.

Q: Why can't the AI fake benchmark numbers?
A: Exact values are inserted into the prompt by code and the verdict uses only verified fields.

Q: Can it compare different fiscal years?
A: It warns about the period mismatch instead of producing a misleading comparison.

Q: Where are the answers stored?
A: chat_messages / benchmark_comparisons tables -- the history page shows them again.